GYR-O PAGES · THE FORMAT

GYR-O Pages ·
A Universal Format for Single-Topic Websites, Generated by AI

WHAT THIS IS

It's a universal page layout and organization scheme for lightweight, single-topic websites.

A typical instance output by the generator: worland-dog-law.online. See every GYR-O site.

This page is not a description of the format. It *is* the format, with itself as the subject — so every claim below is demonstrated by the thing you are reading.

free to host · nothing to hack · works on any device · you own it outright

THE PAGE LAYOUT BELOW IS STANDARD ON EVERY GYR-O SITE

A PLAIN ANSWERS ✓

– What is GYR-O Pages?

Every GYR-O page comes in two parts: Intro and Info. Intro, above the double line, is whatever this subject needs — its name, a picture, a paragraph — and it differs from site to site. **Info**, below the line, is the standard layout, the same on every GYR-O site: three columns of plain answers, sources and technical detail. A reader who has used one already knows where to look on the next.

A way of building a single-topic website so that one topic serves two readers at once: the newcomer who knows nothing about it, and the reader who already knows the subject well. Neither is watered down for the other, and neither has to wade through the other.

Its second idea is about footnotes. Instead of being buried at the foot of a page, tucked behind a link, or scattered wherever they fall, every reference is gathered into one column of its own — centered in the layout, given the same visual weight as the writing, and always in the same place.

They are three columns wherever they will fit as three columns. On a wide screen you see all three at once. On a phone they stack, and you read one at a time. Printed, they become one document. And where one column holds more than any column can — an entire municipal code, say — it is given its own page, as at worland-dog-law.online[6]. Stacked, paged or side by side, it is the same three. Usually it is a single page — no system to log into, no monthly bill, nothing that can be switched off.

It is written down as a formal specification, the GYR-O Pages Format Definition: numbered rules, each marked MUST, SHOULD or MAY in the manner of internet standards, against which any GYR-O site can be checked.

– Who is it for?

Anybody who has to be believed by strangers.

- A small business whose customers want to know it is licensed, insured and real.
- A tradesperson, inspector or consultant whose claims rest on standards and certifications.
- A campaigner or citizen quoting a law, an ordinance, a filing.
- A researcher publishing work that a skeptic should be able to check.

The common thread is not the subject. It is that a reader might reasonably ask *says who?* — and deserves an answer better than a bare link.

– Why three columns rather than one long page?

Because one subject has to be met at more than one depth, and mixing the depths serves nobody. A customer scrolling past technical specifications stops reading. A specialist scrolling past an explanation of basic terms also stops reading.

- **Column A** — the newcomer. What is this, what does it cost, why should I care.
- **Column B** — the checker. Your sources, previewed, so they can verify without leaving.
- **Column C** — the specialist. Full detail, methods and specifications.

Side by side, each picks a lane and stays in it — and can glance sideways when curious.

– Why does the green column come first?

The GYR-O scheme defines standard positions within the page layout, so that essential elements sit in the same place on every GYR-O site. The foremost of those positions, the green column, is

designed for the newcomer. It holds the Plain Answers: the seven to ten most quintessential questions a first-time visitor would ask. The first of them is the first thing anyone reads. Readers who already know the subject can go straight to the red column for the technical detail, so the page serves both kinds of reader at once.

Starting from one provided sentence about the business, AI does deep research, determines the questions, and writes the answers with startling speed. The owner of the site then checks the facts. Without AI, this would necessarily be the work of a long-serving insider, the one person who knows the business well enough to do it but is never free to. The result is what the industry calls a **pillar page**[24], a recognized best practice: one broad topic, with the main questions about it answered on a single page. It also follows Google's published guidance for 2026[25], which rewards original, first-hand answers to the questions people actually ask.

The pillar page is the most under-acknowledged advantage of the format, which is why this answer gives it so much room. A page built this way serves both of its readers at once: the human newcomer, who finds the essential questions answered first, and the search engines and AI systems that read the page on everyone else's behalf, which find one topic covered in full and organized around the questions people actually ask. The GYR-O scheme does not merely support the pillar page; it revolves around it.

– What is the middle column for?

Every source, previewed. When column A or C cites something, the citation is a small red **numeral** [1] — the number you just passed. That is what they are called throughout this page. Clicking one moves you to a card in the middle that tells you what the source is, who made it, what you would find there, and what it is doing in this page.

For a business this column *is* the sales pitch. Not "we are licensed" but a card naming the license, the issuing authority, and a link to the register where anyone can check it in ten seconds. That is a different order of credibility from a badge image.

At the top of the column, before any source, two cards name the most recent changes to the page and a third links the chronicle — the whole record of what has changed and why. A returning reader sees at once what is new. And when a numeral has brought you to a card, a Back button at its foot takes you to the exact spot you left.

– Why the colors?

Because a consistent color code helps readers find their way and recognize what they have seen, and the traffic signal is a code everyone already knows. Green is for the newcomer, yellow for checking the sources, red for the specialist. Orange, which a reader meets only once, at the foot of the page, belongs to the build.

A Harvard and MIT study showed people about 400 published charts and infographics for one second each, then tested which ones they recognized when they came round again. Those with seven or more colors were recognized more often than those with two to six, and both more often than those in a single color or in grays[26]. The study counted colors; it did not test which ones, and nothing in it favors one set of colors over another. So the working premise here is that a color code helps because it is consistent, not because of the particular colors: any code would do, provided it never changes. That is a premise, not a finding. The study also had its limits: it measured whether a picture was recognized again, not how much of it was understood, and after one second rather than a full reading.

Since any code would do, GYR-O borrows the one that needs no teaching. Read left to right, the traffic signal becomes a scale of effort: **green**, go, start here; **yellow**, slow down and check; **red**, stop here only if you mean it. Apple made the same choice at the busiest corner of the computer screen. When it rebuilt its operating system as Mac OS X, the Aqua design it showed at Macworld in January 2000 put red, yellow and green buttons at the upper left of every window[27], "resembling the lights on a traffic signal," as Macworld put it[28]. Borrowed, not invented, and that is the point: nobody has to learn it. GYR-O runs the signal the other way, green first, so every page starts at *go*.

The colors ride only on the page's furniture: column heads, stripes, keys and numerals. The words stay black, and the letters A, B and C and their fixed places carry the same meaning, so a reader who cannot tell the colors apart, or who prints in black and white, loses nothing.

– Why should a reader believe any of it?

They should not have to. That is the whole design.

Transparency is the property this format is built around, and it is not a sentiment — it is a set of rules, each of which *costs the author something*. That cost is the only reason any of them are worth anything.

- **Every claim carries its source, previewed where it is made.** The reader checks without leaving the page, and without taking your word for what the source says.
- **A defect found is published, not quietly fixed.** There is one on this site now: the license arrangement described on Page F gives away more than it should, and it is written up there rather than tidied away first.
- **An argument that would help but cannot be checked is left out.** The Facebook answer drops the most popular argument against feeds — that they punish outbound links — because nobody outside the company can verify it. Weakens the case; keeps it honest.
- **The checks are named, not merely run**, with their count, so the claim can be audited rather than admired.
- **Dead links are marked, not removed.** Deleting a source that died is editing the record.
- **Samples say they are samples**, across the masthead, in a box that cannot be missed.

- **The history is on the page**, in the chronicle, the dated record of every change made to it — including what is still unresolved.

For a business the same property is worth money rather than credit. The tree service sample publishes its federal employer number so that a customer can pull the state's own record of its insurance *without asking the company*. No trade site does that, because it invites checking. Inviting checking is the point.

Nothing here asks to be trusted. It asks to be checked — which is a different and much cheaper thing to offer.

– Why is there no "ibid.", no "see above"?

Because every reference on these pages names *what* it points at, never *where* it sits.

A numeral — one of the small red numbers in the text — points at the card that carries the same number in the middle column. It means the same thing whether you got here by reading from the top, by clicking a numeral in the third column, by searching, or by following somebody else's link straight to the middle of the page. **Ibid.** means "the note immediately before" — and on a page read that way there is no *before*.

It is one property with four consequences, which is why it is worth naming rather than leaving as good manners:

- **Nothing needs abbreviating.** A source appears once, as one card. Citing it three times costs three numerals, not three citations. The problem ibid. was invented to solve does not arise.
- **A machine can check it.** A numeral has a target that either exists or does not. "See above" has no target at all, so no checker can test it and no reader can be sure of it.
- **A copy of the page is self-contained.** The archived capture resolves on its own, because nothing in it depends on the reader's route through it.
- **A quotation survives being cut out.** Lift a paragraph and its references still say what they refer to. Lift one containing *ibid.* and it refers to nothing.

Print conventions — *ibid.*, *op. cit.*, *supra* — are all relative addressing. Each saves a line by assuming a reader moving in one direction through a fixed object. Neither assumption survives a web page, which is presumably why the Chicago Manual of Style stopped recommending *ibid.* at its seventeenth edition, on the ground that it "risks confusing the reader"[23].

The simplicity is a consequence, not the aim. It is what falls out of refusing to let any reference depend on position — and it is the property to protect first if this format is ever changed.

– Why is the link the last thing in each card?

Because when it was first, readers clicked it immediately and never read the preview — which defeated the whole point. That was the single most important discovery in designing this.

By the time you reach the link you already know whether the trip is worth taking. Most readers stop at the preview, satisfied. That is a success, not a failure.

– What happens when one of my links dies?

Links rot. Government pages move, companies fold, archives reorganize. A page that quietly turns into a row of dead ends is worse than no page.

So the links are checked[2] and when one dies it is *disabled, not deleted* — you can still see where it pointed, the preview still says what was there, and a dated note appears at the foot of the middle column saying when the break was found.

Nothing vanishes. A reader years later can see exactly what you cited and that it was live when you cited it.

– What does it cost — and what if the free hosting ends?

Free, without a name of your own. The hosting is free[3], the security certificate is free[4], and an address like `yourname.github.io/site` costs nothing and never lapses. **A little more each year buys your own name instead** — and that is the part you actually own, because a name you rent can be moved anywhere and a free address cannot.

Or share a name. A group can buy one name together — a farm cooperative, say, with `washakiefarms.online` — and every member's site lives under it free, at an address like `eastbench.washakiefarms.online`, with one index page listing them all.

Think of license plates. Everyone gets one; a vanity plate costs a little extra each year, and it is up to you. A name of your own runs roughly $11 a year for a `.com` and up to about $28 a year for a `.online` once its cheap first year is over.

The honest answer to the second half. Nobody has promised free hosting forever, and you should distrust anyone who says otherwise. What protects you is not a promise — it is that your site is a handful of plain files. If the host ever changed its terms you could move the whole thing somewhere else in an afternoon, unchanged, and keep your address. There is nothing to export, nothing to convert, and nothing held hostage.

That is the real answer to lock-in, and it is stronger than "forever" because it does not depend on anyone's goodwill.

– Is it secure?

Yes, and for a better reason than you would expect. It is not that the host has excellent security staff, though it does. **It is that there is nothing there to attack.**

No database, so nothing to inject. No login, so no password to steal. No forms and no server-side code, so nothing to run. No plugins, so nothing to go out of date. No cookies and no analytics, so no visitor data to lose. The entire site is text files being handed out.

The overwhelming majority of small-business sites that get hacked are running a content system with out-of-date add-ons. This has no content system and no add-ons. It cannot go out of date because there is nothing in it that has a version.

– Will Google find it?

Half of that is automatic and half of it is you, and it is worth being clear which is which.

The technical half is as good as it gets by construction[10]. The page is small and loads almost instantly, works on a phone first, uses proper headings and structure rather than a soup of anonymous boxes, and has nothing blocking it from rendering. Those are the things a search engine measures and most sites struggle with.

The other half is your words and who links to you, and no format can supply that. Anybody promising you rankings is selling something.

– Do I really own it?

Mostly yes, and the honest answer has a hole in it that is worth knowing about before somebody else finds it.

You own the domain — you are the named registrant on the public record, and ICANN's own statement of registrant rights includes the right to move the name to a different registrar whenever you like[20]. That portability is what makes an address yours rather than rented.

You own the words, automatically, from the moment you write them. Nothing to file.

You do not own the hosting, and this is the hole. GitHub's terms say plainly that you own your content and that the license you grant them is limited to running and improving the service — and equally plainly that they may "suspend or terminate your access to all or any part of the Website at any time, with or without cause, with or without notice"[19]. Free hosting is free because it is not a promise.

What makes that survivable is the rest of the design: the files are plain files on your own disk and the domain points wherever you tell it, so losing the host costs a Saturday rather than the site. That is the real claim — not that nothing can go wrong, but that nothing that goes wrong takes your work with it.

The full account is on Page F, including a licensing defect in these repositories that is recorded there rather than quietly fixed.

– Why not just keep my Facebook page?

Keep it. This is not an either-or, and anyone who tells you to delete it is wrong. A Facebook page reaches people who are not looking for you; a site of your own answers people who are. They do different jobs.

What a Facebook page cannot do is hold still.

- **A feed is not an address.** Everything there is ordered by when it was posted. The answer you gave in March is not wrong in June, it is just three hundred posts down. A site gives every answer an address that never moves, so "are you insured?" is always in the same place.
- **Machines cannot read it.** Facebook asks crawlers, archivers and link checkers to stay out, and the well-behaved ones obey[16]. Four ordinary fetches of facebook.com while this answer was being written were all refused[17]. Practical effect: nothing can archive it, your own link checker cannot verify it, and another page that wants to cite you as a source cannot quote you.
- **The address is not yours.** It is on somebody else's domain, under their terms, and you cannot move it or take the look of it with you.
- **You keep the copyright and grant away the use of it.** This is the part people get backwards. You do own what you post — on every platform. What you hand over in exchange for posting it is a license: on Facebook a "non-exclusive, transferable, sub-licensable, royalty-free and worldwide" one, covering modification and derivative works[18]. Ownership is not the thing protecting you; the terms of the license are.

So: both. Post on Facebook, where the people already are, and link to the page that stays put. The one you own is the one a stranger can check at nine at night without an account, and the one still there in five years whatever happens to anybody's feed.

One honest correction, because it is the usual version of this argument: it is *not* true that Facebook forbids outbound links. You can post one. Whether the feed shows it to many people is a separate question nobody outside Meta can answer, so this page does not make that claim — the three above are enough, and they are checkable.

– Does it work on a phone?

It is built for the phone first. Below a certain width the three columns become three buttons and you read one at a time; above it, all three sit side by side. Same file, no separate mobile site, nothing to keep in step.

Phone view, the button beside Download PDF, shows this page inside a phone-sized frame at the real widths of common phones and a small tablet — so the phone layout can be checked from a

desktop, without picking up a phone. On a screen that is already phone-sized it never appears, because there is nothing to show you. On this page it is there for every reader, since the layout is part of what the page is demonstrating; on a site built from the format it is for whoever builds the page, and stays hidden from everybody else.

On a long page, once you have scrolled down, a round button with an up arrow appears at the bottom right. It takes you straight back to the top.

It also prints properly, which almost nothing does anymore: all three columns unfold into one document in order, with every collapsed answer opened.

– Do I need GitHub?

Only if you want to make the changes yourself. GitHub is where the files sit, and where the page is served from. Somebody has to put them there once — you, or whoever builds the page for you. A page that goes up and is then left alone needs nothing from anybody after that.

If you do want to change your own page, you never have to see a command line. Two free things and about twenty minutes, once:

1. A **GitHub account**[3] — an email address and a password. Free.
2. **GitHub Desktop**[9], an ordinary application for Mac or Windows. Free. It puts a folder on your computer that is joined to the internet.

A change then goes live like this: edit the file, open GitHub Desktop, type one line saying what you changed, click **Commit**, then click **Push**. Live in about a minute.

Two words to learn. **Commit** means "save this version, with a note." **Push** means "send it to the internet." Everything else in that application can be ignored indefinitely. Column C has the setup written out step by step.

– Does it take a team?

No. One person can write it, publish it and look after it, and that is more of an advantage than it sounds. Every pair of people on a job is one more conversation to keep going: two people have one, five have ten, ten have forty-five[13]. The count climbs like a hockey stick, and column C draws it. Each extra person also brings their own idea of how the thing should work, and every idea has to be squared with every other. One person has none of that — no meetings, no handoffs, no waiting on anybody else before a change can go live.

One is enough, not a limit. Every change is saved with a note saying what changed and who made it, so a second person can join at any time and see how the page got to be the way it is.

– How do I change the words later?

Three ways, and you only need one of them.

1. **Ask whoever built it.** For most subjects that is the whole answer, and it costs the builder a few minutes.
2. **Open the file and type over the words.** The page is one file, and your sentences sit in plain sight between the tags that shape them. The rule is to leave the angle brackets alone and change what is between them.
3. **Ask an AI to make the change.** Say what should be different, in plain English, and read what comes back before it goes up. That is how every site in this family is kept, including this one.

What this format gives up. WordPress[14] exists to give you an editing box: sign in, type, click Update, and the page has changed. That box is the reason for the database behind it, the account you sign in to, the plugins that extend it and the updates that never stop — and for most of what goes wrong on small sites. There is no box here, because there is nothing to sign in to. What you have instead is a file anybody can open in fifty years.

What it asks of you. If you want changes made by asking for them in plain English, you need an AI assistant and enough familiarity with one to ask for what you want and to check what comes back. That is the one skill this format assumes — and teaching it is a main point of the pilot program this format came out of[11]: twelve sites, four workshops open to anyone, and a few local people given a light introduction to AI, so that a subject who wants a change has somebody nearby to ask.

– Who sets one up?

One person, once. Making the repository, putting the files in it, turning the hosting on, buying the domain and pointing it at the host takes about twenty minutes the first time, and column C has it written out step by step. On the sites in this family a script does the domain end of it.

The person the site is *for* need do none of that. On the Washakie County pilot program[11] a subject chooses a name and hands over the material — the ordinance, the price list, the record of work — and signs up for nothing, installs nothing and holds no password. The site is handed over at the end, domain and all.

– Can I publish a running series rather than one page?

Yes, and this is where the format earns its keep for some trades. Each issue is dated, published the day it is written, and *never revised afterward* — superseded by a later issue if it turns out to be wrong, but never quietly edited.

A builder documenting a renovation, an inspector logging visits, a restorer recording a treatment, a researcher recording a day's tests[5]. What you end up with is not marketing. It is a dated, public, source-cited record of what was done and when, which is a materially different thing to own if anyone later disputes it.

Worth saying plainly: if that record matters to you legally, ask a lawyer how it would actually be treated. The format gives you the record; it cannot tell you what it is worth in a dispute.

– What if this all disappears?

Then a copy of it is somewhere else, kept by somebody who is not you.

Every page of every site built this way is sent to the **Internet Archive's Wayback Machine**, which stores its own copy, at its own address, on its own hardware, with the date it was taken[21]. The domain can lapse, the host can change its terms, you can lose interest or die, and the pages still answer.

That is worth something on its own, but four other things come with it:

- **A date you did not set.** The capture shows what the page said and when. If the page later makes a claim about priority — that a finding was published on a particular day — the evidence for it is held by a third party who has no stake in the answer.
- **Proof against quiet editing.** Anyone can compare today's page against a dated capture. That constrains the author, which is the point: it is the same discipline as the chronicle — the page's own dated record of its changes — except that this record cannot be rewritten by the person it is about.
- **A citation that still resolves.** Somebody quoting the page in ten years can point at a copy that does not depend on the author still paying for anything.
- **The sources, too.** With a free account, the capture can be told to take every page the site links to[22]. For this format that matters more than it would elsewhere: column B, the previewed sources, exists so a reader can go and check, and link rot is the one thing that quietly destroys that promise.

What it is not. It is not a backup — it keeps the page as a reader sees it, not the files it was built from. It does nothing whatever for search ranking, and brings no visitors. And it is close to permanent: a capture is public, and removing one means asking the Internet Archive rather than deleting a file.

The reason it belongs in this format rather than being a nicety: every page here tells its reader *do not take my word for it, go and check*. An archived copy is that same sentence turned on the page itself — the one link in the chain the author does not control.

– Can I see some, finished?

Five, and deliberately unalike. Three are real sites doing real work; two are samples, built to show a trade what a page of its own would look like before anybody has to pay for one.

In use:

- baconian-ciphers.online[5] — a research strand published daily, dated and never revised afterward.
- worland-dog-law.online[6] — a town's animal-control ordinance read against its own sources.
- washakie-ai.online[11] — a grant proposal written in the format it proposes to introduce.

Samples, for the trades, the crews and the people who do not have a page yet:

- East Bench Pivot Farm — an invented Washakie County farm: hay by the bale or the ton, beef by the quarter, and the county's real crops, water and rules behind every answer.
- Bighorn Basin Tree Service — an invented tree service, built around a single problem: a stranger standing in their own yard cannot tell whether a contractor is insured. So the insurance is the first answer on the page rather than a line in a brochure, and every part of it links to somebody else's check — the State of Wyoming's for the workers' compensation, the arborists' own register for the credential.
- Cottonwood Draw Farms — an invented farm hiring a seasonal harvest crew, written for the crew rather than for a customer: what the work is, what the law says they are owed, and who to call when something is wrong. It carries a full Spanish peer — la misma página en español — the same answers and the same sources, not a summary.
- Ruby Calder · a US Handbill & Broadsheet page — an invented high school FFA member, as the thing a sixteen-year-old has instead of a résumé: the project, the hours against published thresholds, the events, the office, the service. About a hundred printed cards go with it, with a code on the back and instructions for somebody who has never used one, including what to do when nothing happens.

All four wear SAMPLE across the masthead and carry a dashed box saying plainly that the company, the farm or the student is invented. Underneath they are built exactly as a real one would be, which is the point of showing them.

B SOURCES ✓

HOW TO READ A CARD

Every card below reads the same way, top to bottom. Parts 4 and 5 are the ones a bare footnote never gives you.

1. **Number and kind** — what sort of thing it is
2. **Title** — what it is called
3. **Provenance** — who made it, when, under what license
4. **Preview** — what you would find if you went there
5. **Why it is cited** — what it is doing on this page
6. **The link** — last, and only last

An article, a photograph, a map, a book, a dataset and a film, each carded: see every kind at work →

INTERNAL LINKS

The two most recent content changes, and the record behind them.

28 October 2026 — the formal specification, and the standards it rests on

The format is written down as numbered rules marked MUST, SHOULD or MAY, in the senses fixed by the internet's own RFC 2119 and RFC 8174. Column C now says so, links the Format Definition, and names each outside standard a rule depends on.

28 October 2026 — why the colors

The most visible thing on the page, now explained: a consistent color code helps readers, a Harvard and MIT study suggests any code would do, and GYR-O borrows the traffic signal because nobody has to learn it, as Apple did in 2000.

Chronicle — what has been happening, and why

The whole record in plain language, newest first: what changed, what it was for, and what is still unresolved. Kept so that anyone returning after a few weeks can pick the thread back up without reading a commit history.

EXTERNAL LINKS

Every reference to some other website, previewed. The link is always the last line — read the card first and decide whether you need the round-trip distraction of visiting the original at all.

1 STANDARD · TYPOGRAPHY

Source Serif 4 and IBM Plex Mono

Adobe and IBM · open license · free to use commercially

The two typefaces this format uses: a serif for reading and a monospace for anything a reader might need to type, copy or check. Both are free, both are open-licensed, and both stay legible at small sizes on poor screens.

Why it is cited: as a demonstration of what a source card is for. You now know what these are and why, and almost certainly do not need to visit the page at all.

→ fonts.google.com/specimen/Source+Serif+4

2 METHOD · LINK HEALTH

Why a page cannot check its own links

Cross-origin rules · Mozilla developer documentation

For security reasons a browser will not tell a page the result of a request to somebody else's site. The response comes back opaque, with no status attached, so a healthy page and a dead one look identical from the inside.

Why it is cited: because the honest version matters more than the clever one. Any in-page link checker is guessing, and will mark live sources dead. The check therefore runs on a schedule, away from the reader, and writes a small file that the page reads. Column C has the mechanics.

→ developer.mozilla.org/en-US/docs/Web/HTTP/CORS

3 INFRASTRUCTURE · HOSTING

GitHub Pages

GitHub · static hosting from a repository · free for public repositories

Serves plain files straight from a folder, with no build step and no server to maintain. Publishing is a commit. A custom domain is set by putting the domain name in a one-line file at the top of the folder. Private repositories need a paid plan; public ones are free.

Why it is cited: it is the entire hosting arrangement. It is also the thing you could walk away from in an afternoon if the terms ever changed, which is the point made in column A about lock-in.

→ docs.github.com/en/pages

4 INFRASTRUCTURE · SECURITY

Let's Encrypt

Internet Security Research Group · free automated certificates · non-profit

Issues the certificate that puts the padlock in the address bar, automatically and at no cost, and renews it without being asked. On this arrangement you never touch it: it appears a few minutes after the domain resolves.

Why it is cited: it is the second half of the dollar-a-month figure in column A. Neither the hosting nor the encryption is a cost.

→ letsencrypt.org

5 EXAMPLE · IN USE

A research strand, published daily

Worland, Wyoming · dated issues, never revised after publication

The format carrying heavy technical content: measurements, methods, and the things that failed. Several issues contain corrections to earlier issues, printed as corrections rather than quietly fixed.

Why it is cited: it shows the format under load.

→ baconian-ciphers.online

6 EXAMPLE · IN USE

A municipal ordinance audit

Worland, Wyoming · a citizen reading a town's animal-control law against its sources

The same format with no research content and no daily rhythm at all — a standing site rather than a series, and the page-form variant: its middle column, the ordinance text, holds more than a column can, so it is given a page of its own. That column is almost entirely primary documents, linked to the town's own copies.

Why it is cited: to show the format is not only about science and not about daily publishing. Most sites built this way will be single standing pages.

→ worland-dog-law.online

7 STANDARD · TEXT

CommonMark

A precisely specified version of Markdown

Settles what Markdown actually means — how lists, emphasis and code behave when the informal version is ambiguous. Relevant if you draft in plain text before it becomes a page.

Why it is cited: column C recommends drafting in plain text so your words outlive whatever tool made them. This is the version of plain text worth using.

→ commonmark.org

8 STANDARD · ACCESSIBILITY

WAI-ARIA Authoring Practices — Tabs

W3C Web Accessibility Initiative

The agreed pattern for buttons that switch between panels: which roles to use, what keyboard behavior a reader expects, and how a screen reader should announce a change.

Why it is cited: the A / B / C buttons follow it. Three buttons is a small enough interface that getting it right costs nothing.

→ w3.org/WAI/ARIA/apg/patterns/tabs

9 TOOL · PUBLISHING

GitHub Desktop

GitHub · free application for macOS and Windows · no command line

A window listing the files you have changed, a box to type what changed, and two buttons: Commit, then Push. It keeps a folder on your computer in step with the folder that gets published.

Why it is cited: it is the only software a publisher has to install, and it is the step first-timers expect to be difficult. Download, sign in, choose a folder. That is all of it.

→ desktop.github.com

10 STANDARD · SEARCH

Google Search Essentials

Google · what a search engine actually asks of a page

The published requirements: that a page can be crawled, loads quickly, works on a phone, and uses real headings and structure. It is also explicit that content and reputation, not markup, decide ranking.

Why it is cited: it is the evidence behind the split in column A. The technical requirements are met by construction here; the rest is your words, and no format supplies those.

→ developers.google.com/search/docs/essentials

11 EXAMPLE · IN USE

A grant proposal, written in the format it proposes

Worland, Wyoming · a $10,000 proposal for grass-roots AI in Washakie County

A standing page in column form. Its middle column opens with the two links blocks; a numeral in A or C puts a Back button on the card it opens; a round button returns to the top; and the last thing on the page is a card pointing one level up, to gyro-sites.online.

Why it is cited: most of what column C describes as recent — the Top and Back buttons, the rule keeping numerals out of the header, the one-level-up card — was first asked for there, and then written into the format.

→ washakie-ai.online

12 TOOL · TESTING

Playwright

Microsoft · open source, Apache 2.0 license · drives Chromium, Firefox and WebKit

A library for driving a real browser from a script: open a page at a chosen size, click what a reader would click, and read back what the page did — where things landed, what is visible, whether any script failed.

Why it is cited: it runs the checks described under Quality assurance in column C. Each check is something a reader would notice if it broke.

→ playwright.dev

13 BOOK · SOFTWARE PROJECTS

The Mythical Man-Month

Fred Brooks · Addison-Wesley, 1975; anniversary edition 1995 · linked here to Wikipedia's article on it

Essays drawn from Brooks's years at IBM managing the development of OS/360. The best known is Brooks's law: adding people to a late software project makes it later, partly because every pair of people is a line of communication — $n(n-1)/2$ of them among n people, so fifty developers have 1,225. The book also argues that a system keeps its conceptual integrity only when a single architect, or a very few, decides what goes in and what stays out.

Why it is cited: it is where the count of lines between people comes from, and the case for one mind behind a design. A site one person can keep has no such lines, the point made in column A and drawn in column C.

→ en.wikipedia.org/wiki/The_Mythical_Man-Month

14 TOOL · CONTENT SYSTEM

WordPress

The WordPress project · open source, GPL · the most widely used content system on the web

A publishing platform built around an editing box: sign in, type, click Update, and the page changes. Everything else follows from that box — a database to hold the words, accounts to sign in with, themes and plugins to extend it, and updates to keep all of it safe.

Why it is cited: because "why not WordPress?" is the fair question. It is named for what it does well, so that what this format gives up, and what it therefore never has to maintain, is said plainly rather than implied.

→ wordpress.org

15 REFERENCE · THE HISTORY OF NOTES

Note (typography)

Wikipedia · citing Chuck Zerby, The Devil's Details: A History of Footnotes

What notes are and how they are set. On origins it is careful rather than sweeping: Zerby identifies Richard Jugge's 1568 printing of the Bishops' Bible, where the marks (f) and (g) in Job send the reader to notes at the foot of the page, as an early example. It does not claim a single inventor.

Why it is cited: for the date, and for the limits of the date. The claim in column C is only that sending a reader away from the sentence is a habit four and a half centuries old, not that

anybody invented it on a particular morning.

→ en.wikipedia.org/wiki/Note_(typography)

16 STANDARD · WHAT CRAWLERS ARE TOLD

robots.txt (Robots Exclusion Protocol)

Wikipedia · encyclopedia article · CC BY-SA

A file at the root of a site telling visiting crawlers which parts they may fetch. `User-agent: *` followed by `Disallow: /` asks every robot to stay away from the whole site. The article is explicit that the standard relies on voluntary compliance and is advisory rather than enforceable: malicious robots are unlikely to honor it, and some read it as a list of places to look.

Why it is cited: the Facebook answer in column A turns on what such a file does and does not do, and the distinction between a request and a lock is the part most summaries get wrong.

→ en.wikipedia.org · robots.txt

17 OBSERVATION · MADE FOR THIS PAGE

facebook.com, fetched four times

Tested 24 September 2026 · ordinary web client, no account, no special headers

No preview, and that is the finding. Four URLs were requested: a public page, `/robots.txt`, `/legal/terms`, and the business help center. All four returned `URL is disallowed by robots.txt rules` and no content. A control request to an unrelated site in the same session returned its page normally, so the refusals belong to the destination rather than to the tool. A person in a browser can of course read these pages; an automated client that honors the protocol cannot.

Why it is cited: the claim in column A is that a Facebook page cannot be archived, link-checked or quoted by machine. Rather than assert it, the page tried it and recorded what came back. Anyone can repeat the test in a minute.

→ facebook.com · opens for a person, refused to a client

18 GUIDANCE · WHO OWNS WHAT YOU POST

Intellectual property and social media

European Commission · Executive Agency for Small and Medium-sized Enterprises · 23 January 2020

Users retain copyright on every major platform; what differs is the license granted in return. Facebook takes a “non-exclusive, transferable, sub-licensable, royalty-free, and worldwide” license to host, distribute, modify and create derivative works. Twitter takes a “worldwide, non-exclusive, royalty-free” license usable in any and all media, with sub-licensing. TikTok takes an “unconditional, irrevocable… perpetual worldwide license” extending to third parties and to the creator’s name and likeness.

Why it is cited: it answers the ownership half of the Facebook question with a public institution’s summary rather than a blog’s, and it is here because the platforms’ own terms pages cannot be fetched.

→ intellectual-property-helpdesk.ec.europa.eu

19 TERMS · THE HOST

GitHub Terms of Service

GitHub, Inc. · the agreement covering every site hosted this way

States “You own Your Content”, and scopes the license GitHub takes to storing, hosting, archiving, parsing, displaying and copying it as necessary to provide, develop and improve the service, including training AI features, with an opt-out in account settings. Separately: GitHub may “suspend or terminate your access to all or any part of the Website at any time, with or without cause, with or without notice.” A public repository also grants other users a license to use, display, perform and reproduce it by forking.

Why it is cited: the claim that a subject owns their site is only honest if the one part they do not own is named, in the host’s own words rather than a summary of them.

→ docs.github.com · terms of service

20 POLICY · THE DOMAIN

Registrants’ Benefits and Responsibilities

ICANN · the body that accredits domain registrars

What a domain registrant is entitled to from a registrar: the registration agreement in writing and downloadable, accurate pricing, accessible support and dispute resolution, freedom from deceptive practice, and clear instructions for registering, managing, transferring, renewing and restoring the name. The transfer right is the one that matters: a domain is not tied to the registrar that sold it.

Why it is cited: the domain is the only part of a site that is owned in a way anyone would recognize as ownership, and this is the document that says what that consists of.

→ icann.org · registrants' benefits and responsibilities

21 DOCUMENTATION · WHAT THE ARCHIVING DOES

Save Pages in the Wayback Machine

Internet Archive Help Center · the organization that keeps the copies

Describes capturing a single page as it stands at that moment. States plainly that it "does not save any of the outlinks, and can't be used to initiate a crawl of an entire web site", that submissions are anonymous because the submitter's address is not kept, and that some pages cannot be captured at all because of the site's own security settings or crawling restrictions.

Why it is cited: it is the document that corrected this page. The first description written here said the sources would be captured along with the page; for the plain, no-account form of the tool that is false, and this is where it says so.

→ help.archive.org · save pages in the Wayback Machine

22 ANNOUNCEMENT · THE OPTION THAT MATTERS HERE

The Wayback Machine's Save Page Now is New and Improved

Internet Archive Blogs · 23 October 2019 · still the reference for the options

Sets out what was added to Save Page Now: capturing the requested page together with every page it links to, chosen by a checkbox; producing a screenshot beside the capture; submitting by forwarding an email to savepagenow@archive.org and receiving playback links back; and, for anyone signed in to a free account, keeping the captures in a personal public gallery.

Why it is cited: outlinks are the part of this that this format actually needs, since a page whose argument is "go and check the source" is only as durable as the sources are. This is where that option is documented, and it is the reason an account is worth the five minutes.

→ blog.archive.org · Save Page Now is new and improved

23 STYLE AUTHORITY · THE CONVENTION ABANDONED

What's New in the CMOS 17 Citation Chapters

CMOS Shop Talk · the Chicago Manual of Style's own editors · 4 May 2017

Announcing the seventeenth edition's citation changes. On the abbreviation: "ibid." risks confusing the reader, and so, in a departure from previous editions, CMOS 17 discourages its use in favor of a shortened citation that names the work. It is discouraged rather than forbidden. The same chapters add guidance for citation software, social media and text

messages, video games and interactive books, and permalinks and page numbers in electronic formats.

Why it is cited: the claim in column A is that relative references do not survive non-linear reading. This is the standard-setter conceding the same point for print, and the list of what else it had to add — permalinks, electronic page numbers — shows what forced the concession.

→ cmosshoptalk.com · what's new in the CMOS 17 citation chapters

24 PRACTICE · CONTENT STRATEGY

How to Get Ranked and Read With the Topic Cluster Content Model

Content Marketing Institute · on the pillar-and-cluster model HubSpot introduced in 2017

The pillar page explained: one long page giving a complete overview of a single broad topic, linked to shorter pages on each subtopic. It has been standard search-marketing practice since HubSpot's research introduced it.

Why it is cited: "pillar page" is the industry's term, not Google's. This is where it comes from. A GYR-O site differs in one way: the subtopics are answered on the same page rather than on pages of their own.

→ contentmarketinginstitute.com · the topic cluster content model

25 STANDARD · SEARCH, 2026

A new resource for optimizing for generative AI in Google Search

Google Search Central · May 2026

Google's guide to what its search and its AI answers now reward: unique, non-commodity content with first-hand detail, rather than text that reads like every other site. It also warns against using special markup or question-and-answer formatting as a shortcut. The same month Google stopped showing FAQ rich results in its search pages altogether.

Why it is cited: it is the basis for saying the Plain Answers follow Google's guidance. The value is in answering real questions well, not in the format, which Google no longer displays specially.

→ developers.google.com · a new resource for optimizing for generative AI

26 STUDY · MEMORY AND DESIGN

What Makes a Visualization Memorable?

Borkin, Vo, Bylinskii, Isola, Sunkavalli, Oliva and Pfister · Harvard University and MIT · IEEE, 2013

About 5,700 charts and infographics gathered from scientific papers, news sites, design blogs and government reports. 393 of them were shown to 261 people for one second each, about 87 viewings apiece, and recognition was then tested. Section 7.2 and figure 4 cover color: seven or more colors scored higher than two to six, and higher than one color or a gray scale, with the difference holding after pictograms were taken out.

Why it is cited: it is the evidence that more color helps a picture be recognized. It is not evidence about which colors: the paper counts colors and never compares one palette with another. That is why the page states “any code would do” as a premise rather than a finding.

→ vcg.seas.harvard.edu · the paper, PDF

27 REFERENCE · INTERFACE HISTORY

Aqua (user interface)

Wikipedia

Apple’s Mac OS X interface, introduced at the 2000 Macworld conference in San Francisco and shipped with Mac OS X 10.0 the following year. Its windows carry three buttons on the left of the title bar, red, yellow and green by default, for closing, minimizing and zooming.

Why it is cited: for the date, the place on the screen and the order of the colors.

→ en.wikipedia.org/wiki/Aqua_(user_interface)

28 ARTICLE · MACWORLD, 2001

Watch the Traffic Lights

Henry Bortman · Macworld · 31 May 2001

A tip for people new to Mac OS X, introducing the three window buttons “at the left of each window’s title bar, resembling the lights on a traffic signal.”

Why it is cited: for the traffic-signal comparison, made at the time. Only for that: the same article says yellow zooms and green minimizes, which is the wrong way round. Card 27 gives the right order.

→ macworld.com · Watch the Traffic Lights

29 STANDARD · REQUIREMENT LANGUAGE

Key words for use in RFCs to Indicate Requirement Levels (RFC 2119)

S. Bradner · Internet Engineering Task Force · March 1997 · Best Current Practice 14

A two-page document that fixes the meaning of MUST, MUST NOT, SHOULD, SHOULD NOT and MAY, and their synonyms, as used in internet standards: an absolute requirement, a strong default that may be departed from for good reason, and a truly optional item.

Why it is cited: the Format Definition uses these three words in exactly these senses, so its levels mean what they mean in any internet standard rather than whatever the author intended that day.

→ rfc-editor.org · BCP 14, RFC 2119 and RFC 8174

30 STANDARD · REQUIREMENT LANGUAGE

Ambiguity of Uppercase vs Lowercase in RFC 2119 Key Words (RFC 8174)

B. Leiba · Internet Engineering Task Force · May 2017 · updates RFC 2119, part of Best Current Practice 14

A short update saying the key words carry their special meaning only when written in capitals; in lower case they are ordinary English. It also gives the paragraph a document should include to say it follows these conventions.

Why it is cited: it is why the Definition writes MUST, SHOULD and MAY in capitals, and why a lower-case “should” elsewhere on these pages is not a rule.

→ rfc-editor.org/rfc/rfc8174

31 STANDARD · HTML

HTML Living Standard: the details element

WHATWG · HTML Standard, edition for web developers, section 4.11.1

The browser’s own disclosure widget: a summary line that shows or hides further information or controls, with no script needed.

Why it is cited: every plain answer on a GYR-O page is one of these (rule 2.2), which is why the answers open without scripts and why the browser’s own find can search inside them.

→ html.spec.whatwg.org · the details element

C HOW IT WORKS (DETAILS) ✓

Prior know-how is assumed here, and the items become more technical toward the bottom.

The shape

One HTML file and four small scripts. No framework, no build step, no dependencies except two web fonts, and those fall back to system faces if they fail to load. A typical page is 35–100 KB and has finished rendering before a loading spinner would have appeared. Words are drafted here first in plain text — CommonMark, the precisely specified form of Markdown [7] — so they outlive whatever tool made them.

The scripts are `phone-view.js`, `top-button.js`, `back-link.js` and `search.js`. Each adds one control and what that control needs, nothing more, so a page missing one of them loses that control and still works. Wherever a site carries one, it is the identical file — changed once, in the template new sites are made from, and then copied out — so no copy drifts from the others.

Three columns wide, three buttons narrow, one document printed. The same markup serves all three — nothing is duplicated, because duplicated content is content that will drift.

The formal specification

The format is written down, rule by rule, in the GYR-O Pages Format Definition, version 0.27, a working draft, also kept as a PDF. It is a requirements specification in the style the internet's own standards use. Each rule carries one of three levels, MUST, SHOULD or MAY, in the senses the Internet Engineering Task Force defined in RFC 2119 [29] and narrowed in RFC 8174 to the words in capitals only [30]. Each rule is numbered, so any GYR-O site can be audited against it line by line, and the places where the format's own documents still disagree are listed at the end rather than settled quietly.

Where a rule rests on an outside standard, the standard is named and carded. The words are drafted in CommonMark [7]. Answers fold with the HTML standard's own `<details>` element [31]. The column buttons follow the W3C's WAI-ARIA pattern for tabs [8]. Pages are served over HTTPS with certificates from Let's Encrypt [4], what crawlers may read is set by the Robots Exclusion Protocol [16], and search follows Google Search Essentials [10]. The Definition says what a page must do; it does not claim that every page already does all of it. Where one falls short, that is published as a defect, which is itself rule 6.1.

Why the sources sit in the middle

Position carries meaning. A source column on the far right reads as an appendix; at the foot of the page it reads as an afterthought and goes unread. Between the plain answers and the technical detail it reads as what it is: the common evidence both other columns rest on.

It also means that whichever column you are in, a citation is one glance away rather than one long scroll.

And the card comes to you. Clicking a numeral does not send you off to find something. It brings the card up to where you are looking and marks it for a moment, so the eye lands in the right place instead of searching the column. The movement is not decoration: it is the difference between evidence arriving and evidence having to be hunted for.

Then it brings you back. A Back button appears at the foot of that card, and it returns you to the very numeral you clicked, outlined in gold for a moment, so that reading begins again at the word it stopped on. Both ends of the interruption are carried, not one: the trip out and the trip home are each a single click, and the thread the reader was holding is still in their hand.

That hunt is old. Printed notes at the foot of the page go back at least to 1568[15], and every one of them since has asked the same thing of a reader: leave the sentence, travel to the bottom of the page or the back of the book, find the mark again, come back if you can. Each trip costs a little of the thread the reader was holding. Four and a half centuries of that is long enough to have stopped noticing the cost.

Two blocks, always in the same place

Column B opens with two blocks built exactly alike: a heading, a caption line, then the content. **Internal links**, captioned *The two most recent content changes.*, holds two cards, newest first — each a date and what changed, linked to the place on the page where it changed. Two rather than one because a single line reads as thin, and three begins to look like a log. **External links**, captioned *Every reference to some other website.*, holds the source cards, one for every other site the page mentions.

The headings and captions are word for word the same on every site that carries them, and External always sits below Internal, in one column. A reader who has used one GYR-O site already knows where to look on the next. On a site in page form they head the page that is column B, as on worland-dog-law.online's Ordinance Text page[6], where External links is a plain list.

Citations and the way back

The six parts of a source card are listed at the top of column B, where a reader meets them. Items 4 and 5 do the work. Item 5 is the honest one: it forces the writer to say what the source actually supports, which is where padded citation lists fall apart.

Numerals — the small red numbers that follow a claim, each one a link to the card that supports it — belong in columns A and C, never in the header. A name in the header may still be a plain link.

Arriving at a card from a numeral adds one control at its foot, below the link: a **Back** button, pointing the way you came — `← Back` to column A, `Back →` to column C. It returns you to the very numeral you clicked, outlines it in gold for a moment so the eye finds it, and then goes away. It is a way back, not something the card says; the link is still the last thing the card tells you.

Link health, done honestly

The obvious implementation is wrong. A browser cannot read the status of a cross-origin request: with `mode:'no-cors'` the response is opaque and `status` is 0 whether the page is healthy, missing or on fire[2]. An in-page checker built that way reports false deaths, which is worse than not checking at all.

So the check runs elsewhere, on a schedule, and writes a small file the page reads at load:

```
links.json
{ "checked": "2026-09-10",
  "dead":   { "r4": "2026-08-02" },
  "note":   "one source unreachable since 2 August 2026" }
```

The page fetches that file — same origin, so the status *is* readable — and for anything listed dead it disables the link, strikes it through, and writes the dated note at the foot of column B. If the file is missing or unreadable the page does nothing: no file means no claim, never a false alarm.

On this site, by hand. No scheduler is wired up here yet, so `links.json` carries the date somebody last checked the sources themselves. The page behaves the same either way: it repeats a date a person stands behind, or it says nothing at all.

> **The principle, which generalizes.** A feature that fails silently toward a false positive is worse than an absent feature. This page would rather say nothing about a link than say something untrue about it.

Security: why there is nothing to attack

The usual pitch is that the host has good security people. True, and beside the point. The reason this is hard to compromise is that the attack surface is close to empty:

```
no database          nothing to inject
no login             no credentials to steal
no forms             no input to poison
no server code       nothing to execute
```

```
no plugins          nothing to fall out of date
no cookies          no visitor data held
no analytics        no third party in the page
```

Nearly every small-business site that gets compromised is running a content system with an out-of-date add-on. There is no content system here and no add-ons. Nothing in the page has a version number, so nothing in it can become an old version.

What remains is genuine and worth naming: your account can be phished, and your domain registrar account can be taken over. Both are protected by two-factor authentication, and both are worth turning on before anything else.

Hosting, and what happens if the free tier ends

The pages are plain files served from a folder[3], with the certificate issued automatically[4]. Nobody has promised that will be free forever and the format does not depend on it being so.

The protection against lock-in is structural rather than contractual: there is no database to export, no proprietary format to convert, no theme to rebuild. Copy the folder to any other static host, repoint the domain, and the site is identical. All internal links are relative, so a change of address touches one file.

Search: the half that is automatic

Everything a search engine asks of a page's construction[10] is satisfied here without effort — one small page, no render-blocking scripts, mobile layout first, real headings, real landmarks, descriptive title and meta description, clean relative URLs, no interstitials.

What no format supplies is content worth finding and other sites linking to you. Any claim beyond the technical half should be treated as sales.

One caveat specific to this design: content hidden inside collapsed answers is in the markup and therefore indexable, but a reader arriving from a search lands on a closed accordion. Put the sentence that matters in the summary line, not only inside it.

What a feed cannot do

The question behind this section is the most common objection the format meets: the subject already has a Facebook page, so why build anything? The answer is not that one is better. It is that a feed and an address are different kinds of object, and only one of them can be cited.

Ordering. A feed is sorted by recency, which means the position of a fact decays whether or not the fact does. A price list posted in spring is not retracted in autumn, it is buried, and nothing

distinguishes buried from withdrawn. An address does not decay: `/#hours` is where the hours are, permanently, and updating them is editing one line rather than posting a correction that also sinks.

Their rule, not my tool. Facebook's robots file asks automated clients not to fetch its pages. That is a request rather than a lock — the protocol is advisory and relies on voluntary compliance[16] — but the clients that matter here are exactly the well-behaved ones: search crawlers, web archives, and link checkers like the one described above. Four fetches were attempted from an ordinary web client while this section was written — a public page, the robots file itself, the terms of service, and the business help center — and all four came back `disallowed by robots.txt`. A control fetch of an unrelated site in the same session succeeded, so the refusals are the destination's rule and not a broken tool[17].

That is the concrete loss. A source card in column B previews what is at the other end of a link; it cannot preview a page no client may fetch. So a Facebook page can be linked to, but it cannot be cited in this format — not by choice, but because there is nothing to put in the card.

What is not claimed here. That Facebook demotes posts containing outbound links is widely repeated and poorly evidenced; the only sources are marketing blogs inferring it from their own reach figures. It may well be true. It is not checkable from outside, so it is not used as an argument. The three points above are, which is why they are the ones on the page.

The recommendation that follows is not migration. It is that the feed does what a feed is good at — reaching people who were not looking — and links to the address that holds still.

Devices

A single breakpoint. Above it, three columns with the unselected two dimmed rather than hidden, so the whole page is scannable; a click anywhere in a dimmed column selects it, as its button would. Below it, one column at a time with the buttons acting as tabs. Nothing is a separate mobile site, so nothing can fall out of step. **Phone view**, beside Download PDF, shows the page inside a phone-sized frame at the real widths of common phones and a small tablet, so the phone layout can be checked from a desktop; it never shows on a screen that is already phone-sized. On this page it is there for every reader, because the phone layout is part of what the page demonstrates: its button carries `data-public`. On a site made from the format it is for whoever builds the page: its button stays hidden unless the address carries `?dev`, and then stays on for the rest of that browser tab — the one thing such a page ever stores, and only in the browser of whoever asked for it; `?dev=0` turns it off.

Getting around: the buttons that come and go

Two buttons appear only when they are useful, and neither is printed. Scroll about a screen down and a round button with an up arrow appears at the bottom right. It takes you back to the top — smoothly, or at once if your system asks for reduced motion — and hands the keyboard to the page

title, so the next Tab starts from the top. The page ends with a little blank space, so the last line of the footer can scroll clear of the button on a phone.

The other is the Back button on a source card, described above: it appears when a numeral sends you to column B, and goes away once it has brought you back. Both were first asked for on washakie-ai.online[11] and then written into the format, so every new site starts with them.

Search, with the page as its own index

The box in the navigation bar searches this page and nothing else. When the page loads, `search.js` walks what is actually there — every answer in column A, every card in column B, every section in column C — and builds its index out of the text it finds. Nothing is kept by hand, so the index cannot fall out of step with the page: change a sentence and the next load indexes the sentence you wrote.

It earns its place here for a reason peculiar to this format: the answers in column A are folded shut, and a browser's own find bar will not always look inside a folded answer. A reader could search for a word that is on the page and be told it is not. A result names the column and the heading it came from; choosing one opens that column, unfolds the answer if it is folded, and marks where you landed, the way a numeral does.

No index file, no service, nothing sent anywhere: it runs in the browser, over text the browser already has, and it stores nothing. With scripts off the box never appears, which is the honest state for a control that could not work.

Typography

Body text at 17 px, line height 1.62; on a wide screen each column holds about sixty characters to the line. Figure captions are set at **22 px** — larger than body text, deliberately, because a caption is often the only thing a scanning reader reads; the caption under the picture at the head of this page is one. Three accent colors only: green for structure, gold for the column keys and for marking where you are, red for citations and for whatever a click has just brought you to. A link that has died turns gray and is struck through. Dark mode redefines the nine color properties, and one label color set outside them; the numerals keep their red on a white disc in both themes.

Print and PDF

The button calls the browser's own print dialog. Print rules unhide all three columns, drop the navigation, the floating buttons and the picture at the head of the page, and start each column on a fresh page; just before printing, a script opens every folded answer, and folds them again after — so a printed copy is a linear document in the order A, B, C, and its reader is missing nothing.

Accessibility

The buttons are a real tab set: `role="tablist"`, `aria-selected`, panels labeled by their button, arrow-key navigation[8]. Answers are native `<details>` elements, so the browser's own find searches inside them. With JavaScript disabled entirely the page still reads: all three columns render, on a phone as on a desktop, and every answer still opens with a click. What is lost is whatever moves you around — the switching, the jump from a numeral to its card, the buttons, the striking-through of dead links — and printing then leaves the answers folded.

The Top and Back buttons carry spoken labels, since neither shows much text: *Back to the top of the page,* and *Back to where this source is cited, in column A* (or C).

Ergonomics, interface and experience

Three words for overlapping ideas, and GYR-O answers to each. **Ergonomics**, or human factors, is the oldest and broadest: the science of fitting a tool to the person using it, body and mind, including how much attention it costs. **User interface** (UI) is the visible surface: the boxes, letters, type and color. **User experience** (UX) is the whole encounter over time: whether people find what they came for, at what effort, and whether they come back.

As ergonomics, GYR-O treats the reader's attention as the scarce resource. Visual clutter is cut wherever it can be, and nothing on the page makes the reader hunt: a citation number brings its card to the middle of the screen, a card's preview says what is there before the reader leaves, and a linked passage opens with its section already switched on.

As interface, the controls are few and plain. Two boxes, Intro and Info, decide how much of the page is showing; three letters, A, B and C, name the columns, each with its own color and the same place on every site; and the source card always has the same six parts in the same order. Once learned on one GYR-O site, it is learned for all of them.

As experience, the first-time visitor comes first. Intro is complete on its own, and for many readers it is enough. The first answer in the green column picks up where Intro leaves off, for the visitor who wants more. Regular visitors switch Intro off, and the page remembers that on their own device. Nothing is guessed from behavior and nothing is tracked: how much of the page is on the screen is the reader's choice.

Setting a site up, for someone who has used neither GitHub nor GitHub Desktop

Written for a reader who has never done any of it. No command line at any point.

Once, at the start — about twenty minutes

```
1  Make a free GitHub account.
```

```
2  Make a new repository, named for the site. Check "Public".
   A repository is a folder on their servers that keeps
   every version of everything in it.

3  Install GitHub Desktop and sign in.
   File > Clone repository > pick the one you just made,
   and choose where it sits on your computer. You now have
   an ordinary folder that is joined to GitHub.

4  Put index.html in that folder.

5  On github.com, in the repository: Settings > Pages.
   Source: main branch, root folder. Save.
   The site is live at username.github.io/reponame.

6  OPTIONAL, and the only step that costs money.
   Buy a domain. At the registrar, point it at GitHub:
     four A records on @  ->  185.199.108.153
                              185.199.109.153
                              185.199.110.153
                              185.199.111.153
     one CNAME on www     ->  username.github.io.
   Then add a one-line file called CNAME to the folder
   containing just your domain, and publish it.
   Set the DNS BEFORE publishing that file, or GitHub
   shows a red warning that clears itself once the
   records propagate.

7  Wait for the padlock. A few minutes after the domain
   resolves a certificate is issued automatically and a
   checkbox appears: "Enforce HTTPS". Check it.

8  Turn on two-factor authentication on both the GitHub
   account and the registrar account. These are the only
   two things anyone can actually take from you.
```

Sending a change live — about a minute

```
1  Edit index.html.
2  Open GitHub Desktop. Changes are listed on the left.
3  Type one line in Summary saying what changed.
4  Click "Commit to main".
5  Click "Push origin".
6  Live in about a minute. Add ?v= and today's date to the
   address when checking: the host caches for ten minutes
   and you will otherwise be looking at yesterday.
```

Two words, and no more. **Commit** is "save this version with a note." **Push** is "send it to the internet." Branches, pull requests and merges exist and can be ignored indefinitely.

> **Why this route rather than a website builder.** Because every version is kept, dated and annotated, by default and without effort. For a standing page that is a convenience. For a running record it is the entire point.

> **Keep it small and keep it all.** The slogan belongs to business rather than to software, and it is usually about money: keep the outfit small and the whole of what it earns stays inside it. The arithmetic is the same here, and what is kept is not only money. It is the decisions, one voice from the first line to the last, a record of every change and why it was made, and the hours a larger team would have spent keeping itself in step.

Editing: how a page changes after it goes up

There is no admin panel, and that is deliberate — an admin panel is a login, and a login is the one thing this format does not have. Page D is a form that writes text for you to paste into the file. It holds nothing, protects nothing, and can be thrown away.

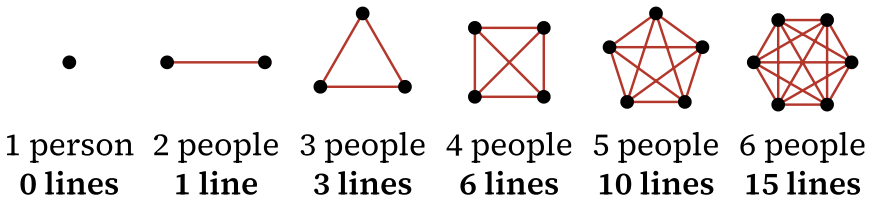
So a change is made to the file, by one of three hands: the person who built the page, the owner in a text editor, or an AI asked in plain English. The third is how every site in this family is kept, and it is the one to be clear-eyed about. What it asks is not programming; it is the ability to say what should be different and to read what comes back. What it costs is a subscription to whichever assistant is used, and the habit of looking at the page before it goes live. On the pilot program this format came out of [11] that help is meant to be local: the workshops are there to leave a few people in the county with a light introduction to AI and the confidence to make a change for a neighbor.

What an editor can break, and how to tell: the words sit between the tags, and the tags are what the browser reads. Delete part of one and the page usually still renders, wrongly — a paragraph swallowed into a heading, a column left half-closed. The checks under *Quality assurance* catch that at both widths; so does opening the page and looking at it.

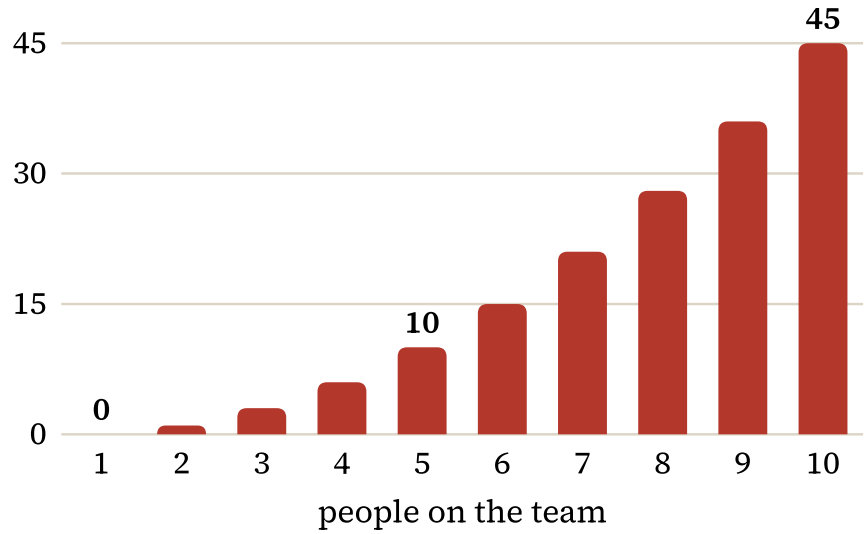
> **Why not WordPress.** WordPress [14] answers this same question with an editing box, and that box is honest work: it is why somebody with no technical knowledge can change their own page without asking anyone. Its price is a database, an account, a theme, plugins and an update every few weeks, each one a thing that can fail or be broken into. This format takes the other side of the trade: nothing to log in to, nothing to update, and a file that will open in any browser for as long as there are browsers.

A team of one

Put a dot for each person at the corners of a regular polygon and join every pair: each line is a conversation the team has to keep going. Three people make a triangle, three lines; four make a square and its two diagonals, six; ten make forty-five. The count is $n(n-1)/2$, which grows roughly with the square of the team, so it climbs slowly at first and then steeply — a hockey stick. Fred Brooks gave the same count in *The Mythical Man-Month* in 1975, as part of the reason that adding people to a late software project makes it later[13].



Lines to keep open, by the size of the team



Every pair of people is one line to keep open.
One person has none; ten have forty-five.

The team this does not need

If the count of lines is the real cost, the test is whether the titles disappear. On a normal web team, a page like this one would be the work of most of these people. Each is a standard job, and each one here is answered by the format rather than by a hire:

```
product manager          the format asks the questions
UX designer              three columns, fixed order
visual designer          type and color, already set
front-end engineer       one file, four small scripts
back-end engineer        there is no server
database administrator   there is no database
DevOps engineer          publishing is a commit
systems administrator    the domain script
security engineer        nothing versioned to patch
```

```
QA engineer             66 checks, and a second reader
release manager         the push is the release
CMS administrator       there is no content system
accessibility engineer  built into the format
SEO specialist          met by construction
analytics engineer      nothing is collected
project manager         nobody to coordinate with
```

Set those sixteen on the drawing above and you have sixteen more corners and a hundred and twenty lines between them, before anybody has written a word for the page. That is the saving, and it is not that the work goes faster: it is that the conversations never have to happen.

Two of those titles are near-twins and both are real jobs: a *product* manager decides what goes on the page, a *project* manager keeps a team in step. The format answers the first. A team of one answers the second.

Two honest limits. This holds for a page of this kind: one subject, a few thousand words, nothing bought or logged into. It does not hold for software in general, and a hospital's scheduling system needs every one of those people. And three jobs do not go away, because they are the work itself — reading the sources, writing the plain answer, and checking that the two match. The format does not remove that labor. It removes everything standing between it and the page.

The talk is only the visible part of the cost. Brooks's larger point was conceptual integrity: a design holds together when a single architect, or a very few, decides what goes in and what stays out, and every added mind is another set of assumptions to reconcile. A site this size needs one person: no meetings, no handoffs, no one else's conventions to learn. Every site in this family is built that way, by one person working with an AI, which drafts and checks while the person decides and publishes. Between people, the count is zero.

One is the fewest it needs, not the most it allows. The site lives in git, which keeps every change with its author, its date and its note, so a second person can take up the same site at any time and read how it got the way it is. If two people change the same lines at once, git stops and shows both versions, and one of them chooses.

The cost of one is the other side of the same coin: with one keeper, what that person knows can leave with them. The format answers that as far as a format can. The page explains its own workings, here in column C; git keeps every change with its note; and nothing in the page's machinery has a version to go stale while nobody is looking after it. The words can still go out of date, and links still die — the link check above says when one has.

A running record

Optional, and not what most sites will do. Each issue is dated, published the day it is written, and superseded rather than edited — a later issue may correct an earlier one, but the earlier one stays exactly as published.

Files go `/index.html` for the current issue and `/YYYY-MM-DD/index.html` for the permanent dated copy, so a link sent today still resolves next year. The version history underneath is a second, independent record of when each word was written.

A site with more than one series gives each a folder of its own, as baconian-ciphers.online does with `/biliteral/` and `/word-cipher/`[5]. The dated copies are never touched again — not even to add a button the current copy has since gained.

A family of sites

Sites built this way can form a family, wired upward. On an instance that sits under another site in this family, the last thing on the page, below the footer, is a card headed ↑ **One level up**: the name of the site above it, and one line saying what that site is, in that site's own words. worland-dog-law.online points to washakie-ai.online[11], and washakie-ai.online points here. This site is the class at the top, so it carries no such card; baconian-ciphers.online carries none either, because it links upward to a site outside this family. The map at the head of the page shows the four sites and which way their links point.

An instance adds what its own subject needs. worland-dog-law.online, three pages long, widens the same search box to cover all three, so a reader on any page finds a section on another[6]; baconian-ciphers.online, a running record in two series, lists every issue of both on its front page[5]. What proves itself on one instance is written back into the format — here, and in the template new sites are made from — so the next site starts with it.

Page D — the setup sheet nobody sees

A fourth page exists, unlinked and marked not to be indexed. It is the setup sheet: the header, the picture at the head of the page, the most-recent-change card, the source cards, the one-level-up card, and the last link-check result. Its header knows the two kinds of page — a standing page and a running series — and each block is written out as text to paste into the file.

It is not an admin panel and has no login, because there is nothing behind it to protect. It is a form that produces text you paste into a file. It is kept out of the navigation because it belongs to whoever builds the site, not to whoever reads it. Somebody who came to read about a town's dog ordinance has no use for a form that writes HTML, and putting it in the menu would only invite the question of what it is for.

Quality assurance

Changes are checked in two different ways before they are handed over for publishing, because each catches what the other misses.

By machine. A script drives a real browser[12] over this page and its template, at desktop width (1440 px) and phone width (390 px), in light and dark, and makes 68 checks: no sideways scrolling and no script errors at either width; the most-recent-change card, which must open the passage it names; numerals, which must reach their cards and come back through the Back button, from column A and from column C; the Top button; Phone view, shown to every reader here and hidden on the template until the address carries `?dev`; what printing leaves out, and that every answer opens for print and folds again after; all three columns with JavaScript off, on a phone and on a desktop; and every block Page D writes. A change is not handed over until all 68 pass. The script is kept with the format's tools, so the checks can be run again.

By a second reader. A script checks only what someone thought to test. So larger changes are also read, claim by claim, against the code and against the sites the text describes, by a reviewer that had no part in writing them — a second AI agent, given the files and told to be skeptical. Its first reading of the 21 September update found seventeen problems, several in older text: the page said folded answers printed, and they did not; it said lines were capped near 62 characters, and on a tablet they ran past a hundred; it listed as one of this page's failures something that could not be made to happen here. Each was fixed or reworded; a second reading confirmed the fixes and caught three more, since fixed.

The links, by hand. Whether each source is still there is checked separately, as *Link health, done honestly* describes: every source is fetched, and `links.json` records the date.

> **The principle.** A test checks what a page does; only a reader can check what it says. This page claims that every claim on it is demonstrated by the page itself, and the second reading is how that claim is kept honest.

Sending the pages somewhere the author cannot reach

A command on the author's own machine walks the volume where the sites live, finds each one the way the link checker does — a folder holding both a `CNAME` and an `index.html` — and turns every page it finds into the address that page is served at, using the directory form where the file is an `index.html`. At the time of writing that is **32 pages across four sites.** New sites are picked up without anyone editing a list.

Three kinds of file are deliberately left out. Superseded drafts and anything waiting to be deleted, because archiving those makes a permanent public record of work that was withdrawn. And the `claude/` folder, which by convention holds session notes rather than site content — those are published as a side effect of publishing everything else, and there is no case for making them permanent.

Anything captured in the last thirty days is skipped. The Wayback Machine is asked what it already has before anything is sent, so running the command often costs almost nothing and a page is not

captured twice in a week for no reason.

Success is measured by outcome, not by the reply. Submitting a page returns a response saying the request was accepted, which is not the same as a capture existing. So after everything has been sent, the command waits and then asks the Wayback Machine, page by page, whether a capture dated today is actually there, and reports only what it can confirm. Anything unconfirmed is named rather than glossed; captures can take minutes, and the next run simply sends it again.

Two modes, and the difference is an account. Without one, each page is submitted on its own and the pages are all that is kept. With a free archive.org account, two keys go in a file beside the command, and the site's front page can be submitted with outlinks — which captures the sources cited in column B along with the page citing them[22]. The command works either way and says which mode it is in before it starts. If the archive begins refusing requests, the wait between pages doubles rather than the run failing.

Nothing is committed, pushed or altered on disk; the only thing written is a dated line in a log beside the command. And because a capture is public and effectively permanent[21], the command lists exactly what it is about to send and asks once before sending any of it.

© 2026 Walter Wilkinson, creator of the GYR-O website format. The underlying template is free for reuse, under the MIT license. View details: license and copyright